

**Name: RahimaSiddiqui(2303.KHI.DEG.030)**

**Peer name: Osama Abdul Razzak (2303.KHI.DEG.029)**

**Peer name:M Humza Moeen(2303.KHI.DEG.019)**

## Assignment 5.3

Read data from source to DataFrame in local Spark setup and display DataFrame schema.

tasks/5\_data\_pipelines/day\_3\_spark/data\_assignment

For numerical columns, calculate minimum, maximum and average values.

For categorical columns, create and apply UDF that will change the last letter of every word to "l".

Sort DataFrame by the first column and save the results to the Parquet file.

First, we import these libraries that we are used in this assignment.

```
[16]: from time import sleep

import pyspark
from IPython.display import clear_output, display
from pyspark.sql import SparkSession
from pyspark.sql.functions import col, from_json, udf, when, split
from pyspark.sql.streaming import StreamingQuery
from pyspark.sql.types import DateType, IntegerType, StringType, StructType
```

Then we import dataset and naming the columns according to dataset.

```
[17]: spark = SparkSession.builder.appName("Assignment pyspark").getOrCreate()
titanic = spark.read.option("header", "true").option("inferSchema", "true").option("header", "false").csv("workspace/data/titanic.csv")
columns = ["PassengerId", "Survived", "Pclass", "Name", "Sex", "Age", "SibSp", "Parch", "Ticket", "Fare", "Cabin", "Embarked", "Timestamp"]

titanic = titanic.toDF(*columns)
titanic.printSchema

[17]: <bound method DataFrame.printSchema of DataFrame[PassengerId: int, Survived: int, Pclass: int, Name: string, Sex: string, Age: int, SibSp: int, Parch: int, Ticket: string, Fare: double, Cabin: string, Embarked: string, Timestamp: timestamp]>
```

Then we modifying the schema and changing the data type of survived to string and converting the binary values into string so that we can categorize it.

```
[18]: titanic = spark.read.option("header", "true").option("inferSchema", "true").csv("workspace/data/titanic.csv")
columns = ["PassengerId", "Survived", "Pclass", "Name", "Sex", "Age", "SibSp", "Parch", "Ticket", "Fare", "Cabin", "Embarked", "Timestamp"]

titanic = titanic.toDF(*columns)
titanic = titanic.withColumn("Survived", when(col("Survived") == 0, "No").otherwise("Yes").cast("string"))
titanic.show()
```

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked	Timestamp
1	Yes	1	Cumings, Mrs. Joh...	female	38	1	0	PC 17599	71.2833	C85	C	2020-01-01 13:44:48
3	Yes	3	Heikkinen, Miss. ...	female	26	0	0	STON/O2. 3101282	7.925	null	S	2020-01-01 13:38:11
4	Yes	1	Futrelle, Mrs. Ja...	female	35	1	0	113803	53.1	C123	S	2020-01-01 13:32:00
5	No	3	Allen, Mr. Willia...	male	35	0	0	373450	8.05	null	S	2020-01-01 13:36:30
6	No	3	Moran, Mr. James	male	null	0	0	330877	8.4583	null	Q	2020-01-01 13:31:39
7	No	1	McCarthy, Mr. Tim...	male	54	0	0	17463	51.8625	E46	S	2020-01-01 13:37:31

## This is our updated schema

```
titanic.printSchema()

root
 |-- PassengerId: integer (nullable = true)
 |-- Survived: string (nullable = false)
 |-- Pclass: integer (nullable = true)
 |-- Name: string (nullable = true)
 |-- Sex: string (nullable = true)
 |-- Age: integer (nullable = true)
 |-- SibSp: integer (nullable = true)
 |-- Parch: integer (nullable = true)
 |-- Ticket: string (nullable = true)
 |-- Fare: double (nullable = true)
 |-- Cabin: string (nullable = true)
 |-- Embarked: string (nullable = true)
 |-- Timestamp: timestamp (nullable = true)
```

```
titanic.show()
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+
|PassengerId|Survived|Pclass|          Name|  Sex| Age|SibSp|Parch|          Ticket|  Fare|Cabin|Embarked|          Times
tamp|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+
|          2|    Yes|    1|Cumings, Mrs. Joh...|female| 38|    1|    0|          PC 17599|71.2833|  C85|          C|2020-01-01 13:4
4:48|
|          3|    Yes|    3|Heikkinen, Miss. ...|female| 26|    0|    0|STON/O2. 3101282|  7.925| null|          S|2020-01-01 13:3
8:11|
|          4|    Yes|    1|Futrelle, Mrs. Ja...|female| 35|    1|    0|          113803|  53.1| C123|          S|2020-01-01 13:3
2:00|
|          5|    No|    3|Allen, Mr. Willia...|  male| 35|    0|    0|          373450|   8.05| null|          S|2020-01-01 13:3
6:30|
|          6|    No|    3|    Moran, Mr. James|  male|null|    0|    0|          330877|  8.4583| null|          Q|2020-01-01 13:3
1:39|
```

Then we are checking the minimum, maximum and average value of the numerical data

```
categoricals_columns = ["Sex", "Cabin", "Embarked", "Survived"]

def change_last_letter_after_space(word):
    if word is not None:
        words = word.split()
        for i in range(len(words)):
            words[i] = words[i][:-1] + "1"
        return " ".join(words)
    return word

change_last_letter_udf = udf(change_last_letter_after_space, StringType())
for column in categorical_columns:
    titanic = titanic.withColumn(column, change_last_letter_udf(titanic[column]))
titanic.show()
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|PassengerId|Survived|Pclass|      Name|  Sex| Age|SibSp|Parch|      Ticket|  Fare|Cabin|Embarked|      Times
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|      2|      1|     1|Cumings, Mrs. Joh...|female| 38|  1|  0|      PC 17599|71.2833| C81|      1|2020-01-01 13:4
4:48|
|      3|      1|     3|Heikkinen, Miss. ...|female| 26|  0|  0|STON/O2. 3101282|  7.925| null|      1|2020-01-01 13:3
8:11|
|      4|      1|     1|Futrelle, Mrs. Ja...|female| 35|  1|  0|      113803|  53.1| C121|      1|2020-01-01 13:3
2:00|
|      5|      0|     3|Allen, Mr. Willia...|male| 35|  0|  0|      373450|   8.05| null|      1|2020-01-01 13:3
6:30|
|      6|      0|     3|      Moran, Mr. James|male|null|  0|  0|      330877|  8.4583| null|      1|2020-01-01 13:3
1:39|
```

Then we apply udf function to change the last letter of categorical data to 1 and here is the result.

```
categoricals_columns = ["Sex", "Cabin", "Embarked", "Survived"]

def change_last_letter_after_space(word):
    if word is not None:
        words = word.split()
        for i in range(len(words)):
            words[i] = words[i][:-1] + "1"
        return " ".join(words)
    return word

change_last_letter_udf = udf(change_last_letter_after_space, StringType())
for column in categoricals_columns:
    titanic = titanic.withColumn(column, change_last_letter_udf(titanic[column]))
titanic.show()
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked	Times
4:48	2	Ye1	1	Cumings, Mrs. Joh...	femal1	38	1	0	PC 17599	71.2833	C81	1	2020-01-01 13:4
8:11	3	Ye1	3	Heikkinen, Miss. ...	femal1	26	0	0	STON/O2. 3101282	7.925	null	1	2020-01-01 13:3
2:00	4	Ye1	1	Futrelle, Mrs. Ja...	femal1	35	1	0	113803	53.1	C121	1	2020-01-01 13:3
6:30	5	N1	3	Allen, Mr. Willia...	mal1	35	0	0	373450	8.05	null	1	2020-01-01 13:3
1:39	6	N1	3	Moran, Mr. James	mal1	null	0	0	330877	8.4583	null	1	2020-01-01 13:3

Then we sorting the data by its first column

```
sorted_data = titanic.orderBy(titanic.columns[0])
sorted_data.show()
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked	Times
4:48	2	Ye1	1	Cumings, Mrs. Joh...	femal1	38	1	0	PC 17599	71.2833	C81	1	2020-01-01 13:4
8:11	3	Ye1	3	Heikkinen, Miss. ...	femal1	26	0	0	STON/O2. 3101282	7.925	null	1	2020-01-01 13:3
2:00	4	Ye1	1	Futrelle, Mrs. Ja...	femal1	35	1	0	113803	53.1	C121	1	2020-01-01 13:3
6:30	5	N1	3	Allen, Mr. Willia...	mal1	35	0	0	373450	8.05	null	1	2020-01-01 13:3
1:39	6	N1	3	Moran, Mr. James	mal1	null	0	0	330877	8.4583	null	1	2020-01-01 13:3
7:31	7	N1	1	McCarthy, Mr. Tim...	mal1	54	0	0	17463	51.8625	E41	1	2020-01-01 13:3
9:08	8	N1	3	Palsson, Master. ...	mal1	2	3	1	349909	21.075	null	1	2020-01-01 13:4
3:42	9	Ye1	3	Johnson, Mrs. Osc...	femal1	27	0	2	347742	11.1333	null	1	2020-01-01 13:3
7:52	10	Ye1	2	Nasser, Mrs. Nich...	femal1	14	1	0	237736	30.0708	null	1	2020-01-01 13:3

Then we saving the resultant dataset in parquet.

```
] : try:
    sorted_data.write.mode('overwrite').parquet("titanic_results.parquet")
except:
    print('Expection caught')
```